

Assignment 1
CSE471 : Statistical Methods in Artificial
Intelligence (SMAI)
Spring 2018
Submission Deadline : 23.55 Hrs. 23/02/2018

General Instructions

1. Assignment can be implemented in Matlab, Python or R.
2. Ensure that submitted assignment is your original work. Please do not copy any part from any source including your friends, seniors and/or the internet. If any such attempt is caught then serious actions including an F grade in the course is possible.
3. A single .pdf file needs to be uploaded to the Courses Portal. The file should contain your answers as well as the code you have written and its output. Include the assignment number, your name and roll number at the top-left of the first page of your submission.
4. Your grade will depend on the correctness of answers and output. In addition, due consideration will be given to the clarity and details of your answers and the legibility and structure of your code.

Problem 1 : Voted Perceptron [10 Marks]

The Perceptron algorithm converges only when the data is linearly separable. It does not converge when the data is non-separable. Also, the performance of Perceptron on unseen test data is not that amazing. In order to make it more competitive with other learning algorithms, it needs to be modified a bit to get better generalization.

Consider a data set with 10,000 examples. Suppose that after the first 100 examples, the Perceptron has learned a really good classifier. Its so good that it goes over the next 9899 examples without making any updates. It reaches the 10,000th example and makes an error. It updates. For all we know, the

update on this 10,000th example completely ruins the weight vector that has done so well on 99.99%. What we would like is for weight vectors that survive a long time to get more say than weight vectors that are overthrown quickly. One way to achieve this is by voting. As the perceptron learns, it remembers how long each hyperplane survives. At test time, each hyperplane encountered during training votes on the class of a test example. If a particular hyperplane survived for 20 examples, then it gets a vote of 20. If it only survived for one example, it only gets a vote of 1. In particular, let $(\mathbf{w}_1, b_1), \dots, (\mathbf{w}_K, b_K)$ be the K weight vectors encountered during training, and c_1, \dots, c_K be the survival times for each of these weight vectors. (A weight vector that gets immediately updated gets $c = 1$; one that survives another round gets $c = 2$ and so on.) Then the prediction on a test point is:

$$\hat{y} = \text{sign}\left\{\sum_{k=1}^K c_k \text{sign}(\mathbf{w}_k^T \mathbf{x} + b_k)\right\}$$

This algorithm is known as **voted Perceptron** works quite well in practice, and there is some nice theory showing that it is guaranteed to generalize better than the vanilla Perceptron.

Algorithm 1: Voted Perceptron

Input: $\mathcal{S} = \{(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_m, y_m)\}$

Output: $(\mathbf{w}_1, b_1, c_1), \dots, (\mathbf{w}_n, b_n, c_n)$

begin

Initialization $n = 1, \mathbf{w}_1 = \mathbf{0}, b_1 = 0, c_1 = 1$;

for $iter \leftarrow 1$ **to** Num_Epochs **do**

for $i \leftarrow 1$ **to** m **do**

if $y_i(\mathbf{w}_n^T \mathbf{x}_i + b_n) \leq 0$ **then**

$n = n + 1$;

$\mathbf{w}_n = \mathbf{w}_{n-1} + y_i \mathbf{x}_i$;

$b_n = b_{n-1} + y_i$;

$c_n = 1$

end

else

$c_n = c_n + 1$

end

end

end

end

Assignment Problem

1. Write the program to implement Voted Perceptron.
2. Consider the following UCI Datasets for classification

- Breast Cancer Dataset: <http://archive.ics.uci.edu/ml/machine-learning-databases/breast-cancer-wisconsin/>
- Ionosphere Dataset: <http://archive.ics.uci.edu/ml/machine-learning-databases/ionosphere/>

Remove the rows corresponding to missing values in Breast Cancer Dataset.

3. Run the vanilla Perceptron as well as Voted Perceptron on both the datasets.
4. Report the following:
 - (a) Take different values for number of Epochs, say 10,15, 20, 25, 30, 35, 40, 45, 50. One Epoch is a complete pass of the training set from the algorithm.
 - (b) For each Epoch value in the above, report 10-fold cross validation accuracies of both Perceptron and Voted Perceptron (for both the datasets). (For those who don't know cross validation accuracy, can read the link below: http://scikit-learn.org/stable/modules/cross_validation.html)
 - (c) Plot the number of Epochs versus Cross validation accuracy for both approaches and for both the datasets.
 - (d) Comment on the performance of two approaches. Give reasons.

Problem 2: Least Square Approach [5 Marks]

samples	C_1		C_2	
	x_1	x_2	x_1	x_2
1	3	3	-1	1
2	3	0	0	0
3	2	1	-1	-1
4	0	2	1	0

Table 1: Dataset 1 for Problem 3

Table 1 and Table 2 show two different datasets for binary classification problem. Do the following:

1. Write the program to find the linear classifier using least square approach.
2. Write the program to find the linear classifier using Fisher's linear discriminant.

samples	C_1		C_2	
	x_1	x_2	x_1	x_2
1	3	3	-1	1
2	3	0	0	0
3	2	1	-1	-1
4	0	1.5	1	0

Table 2: Dataset 2 for Problem 3

3. Plot the data points in Table 1 with different colors for different classes. Find the classifiers learnt using both least square approach as well as Fisher's linear discriminant analysis. Plot both the classifiers on the same graph.
4. Plot the data points in Table 2 with different colors for different classes. Find the classifiers learnt using both least square approach as well as Fisher's linear discriminant analysis. Plot both the classifiers on the same graph.
5. Comment on the difference in the classifier learnt in the above two cases. Give reasons.

Problem 3: Latent Semantic Analysis [10 Marks]

In this particular task, you are given a dataset of articles, each belonging to one of the 5 categories. Convert the documents to word-frequency vectors. Finally form Tf-idf matrix. Apply dimensionality reduction using SVD to decrease the number of features with minimal information loss. That is, find the new feature vectors for each document by projecting the original tf-idf vector to the space spanned by the singular vectors. This way, we will get new representation for every document/word/query. Now, the tasks involved are

1. Divide the whole dataset into training and testing sets. That is, for each class, keep 80% of its documents in training set and the remaining 20% in the test set.
2. Approach 1: Learn a multi-class classifier using Perceptron. This can be done by using one versus rest approach. Now, given a set of test documents, classify each of them into one of the 5 classes using One Vs All Voted Perceptron (Take help from TAs to understand this). Report the accuracy on training and test datasets.
3. Approach 2: Given a test document, retrieve 10 most relevant documents (based on cosine similarity) to the given document. Examine their labels. Assign the most occurring label in those 10 as the label of the test document. Report the accuracy of this approach.

4. Also Make a report of the change in accuracies for different levels of dimensionality reduction (depending upon how many singular vectors are considered) and use the optimal number of features in the final submissions above.